

✓ Gaussian process in a fully Bayesian mode

This notebook demonstrates an application of the Gaussian process class in GPax package to a 1D problem.

Prepared by Maxim Ziatdinov (2024)

Edited by Josiah Park for the present purposes (2024)

Enable browser
Settings to
executions

OK

Gaussian process (GP) is a powerful tool for reconstructing with quantified uncertainty an unknown ("black-box") function over a low-dimensional parameter space from sparse measurements. Formally, given a set of observed data points $(x_i, y_i)_{i=1, \dots, N}$ and assuming normally distributed observation noise ε , the GP aims to reconstruct the function $f(x)$ such as $y_i = f(x_i) + \varepsilon_i$, with f sampled from a multivariate normal prior distribution, $f \sim \text{MVN}(0; K(x_i, x_j))$. The functional form of the kernel K is chosen prior to the experiment, and its hyperparameters are inferred from the observations using either Markov chain Monte Carlo methods or stochastic variational inference.

✓ Install & Import

Install GPax package:

```
!pip install -q gpax
```



```
===== 103.4/103.4 kB 2.9 MB/s eta 0:00:00
===== 371.7/371.7 kB 6.3 MB/s eta 0:00:00
===== 351.1/351.1 kB 6.0 MB/s eta 0:00:00
```

```
!pip3 install openmeteo-requests
!pip3 install requests-cache retry-requests numpy pandas
import openmeteo_requests
```

```
import requests_cache
import pandas as pd
from retry_requests import retry
```

```
# Setup the Open-Meteo API client with cache and retry on error
cache_session = requests_cache.CachedSession('.cache', expire_after = 3600)
retry_session = retry(cache_session, retries = 5, backoff_factor = 0.2)
openmeteo = openmeteo_requests.Client(session = retry_session)
```

```
# Make sure all required weather variables are listed here
# The order of variables in hourly or daily is important to assign them correctly below
url = "https://api.open-meteo.com/v1/forecast"
params = {
    "latitude": 37.8715,
    "longitude": 122.273,
    "current": "surface_pressure",
    "hourly": "surface_pressure",
    "start_date": "2024-07-22",
    "end_date": "2024-10-14"
}
responses = openmeteo.weather_api(url, params=params)
```

```
# Process first location. Add a for-loop for multiple locations or weather models
response = responses[0]
print(f"Coordinates {response.Latitude()}°N {response.Longitude()}°E")
print(f"Elevation {response.Elevation()} m asl")
print(f"Timezone {response.Timezone()} {response.TimezoneAbbreviation()}")
print(f"Timezone difference to GMT+0 {response.UtcOffsetSeconds()} s")
```

```
# Current values. The order of variables needs to be the same as requested.
current = response.Current()
current_surface_pressure = current.Variables(0).Value()
```

```
print(f"Current time {current.Time()}")
print(f"Current surface_pressure {current_surface_pressure}")
```

```
# Process hourly data. The order of variables needs to be the same as requested.
hourly = response.Hourly()
hourly_surface_pressure = hourly.Variables(0).ValuesAsNumpy()
```

```

hourly_data = {"date": pd.date_range(
    start = pd.to_datetime(hourly.Time(), unit = "s", utc = True),
    end = pd.to_datetime(hourly.TimeEnd(), unit = "s", utc = True),
    freq = pd.Timedelta(seconds = hourly.Interval()),
    inclusive = "left"
)}
hourly_data["surface_pressure"] = hourly_surface_pressure

hourly_dataframe = pd.DataFrame(data = hourly_data)
print(hourly_dataframe)
hourly_dataframe.shape

```

Enable browser
Settings to
executions

```

Requirement already satisfied: openmeteo-requests in /usr/local/lib/python3.10/dist-packages (1.3.0)
Requirement already satisfied: openmeteo-sdk>=1.4.0 in /usr/local/lib/python3.10/dist-packages (from openmeteo-requests) (1.17.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from openmeteo-requests) (2.32.3)
Requirement already satisfied: flatbuffers>=24.0.0 in /usr/local/lib/python3.10/dist-packages (from openmeteo-sdk>=1.4.0->openmeteo-requests) (24.0.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->openmeteo-requests) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->openmeteo-requests) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->openmeteo-requests) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->openmeteo-requests) (2024.8.1)
Requirement already satisfied: requests-cache in /usr/local/lib/python3.10/dist-packages (from requests->openmeteo-requests) (1.2.1)
Requirement already satisfied: retry-requests in /usr/local/lib/python3.10/dist-packages (2.0.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.26.4)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.2.2)
Requirement already satisfied: attrs>=21.2 in /usr/local/lib/python3.10/dist-packages (from requests-cache) (24.2.0)
Requirement already satisfied: cattrs>=22.2 in /usr/local/lib/python3.10/dist-packages (from requests-cache) (24.1.2)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests-cache) (4.3.6)
Requirement already satisfied: requests>=2.22 in /usr/local/lib/python3.10/dist-packages (from requests-cache) (2.32.3)
Requirement already satisfied: url-normalize>=1.4 in /usr/local/lib/python3.10/dist-packages (from requests-cache) (1.4.3)
Requirement already satisfied: urllib3>=1.25.5 in /usr/local/lib/python3.10/dist-packages (from requests-cache) (2.2.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: exceptiongroup>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from cattrs>=22.2->requests-cache) (1.2.1)
Requirement already satisfied: typing-extensions!=4.6.3,>=4.1.0 in /usr/local/lib/python3.10/dist-packages (from cattrs>=22.2->requests-cache) (4.12.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22->requests-cache) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22->requests-cache) (3.10)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22->requests-cache) (2024.8.1)
Coordinates 37.875°N 122.25°E
Elevation 0.0 m asl
Timezone None None
Timezone difference to GMT+0 0 s
Current time 1728955800
Current surface_pressure 1019.2000122070312

```

	date	surface_pressure
0	2024-07-22 00:00:00+00:00	1001.200012
1	2024-07-22 01:00:00+00:00	1001.400024
2	2024-07-22 02:00:00+00:00	1000.599976
3	2024-07-22 03:00:00+00:00	1000.000000
4	2024-07-22 04:00:00+00:00	999.599976
...
2035	2024-10-14 19:00:00+00:00	1016.099976
2036	2024-10-14 20:00:00+00:00	1016.099976
2037	2024-10-14 21:00:00+00:00	1016.400024
2038	2024-10-14 22:00:00+00:00	1017.000000
2039	2024-10-14 23:00:00+00:00	1018.000000

```

[2040 rows x 2 columns]
(2040, 2)

```

Import needed packages:

```

!pip install matplotlib
import matplotlib.pyplot as plt

# Plotting the data
plt.figure(figsize=(14, 6))

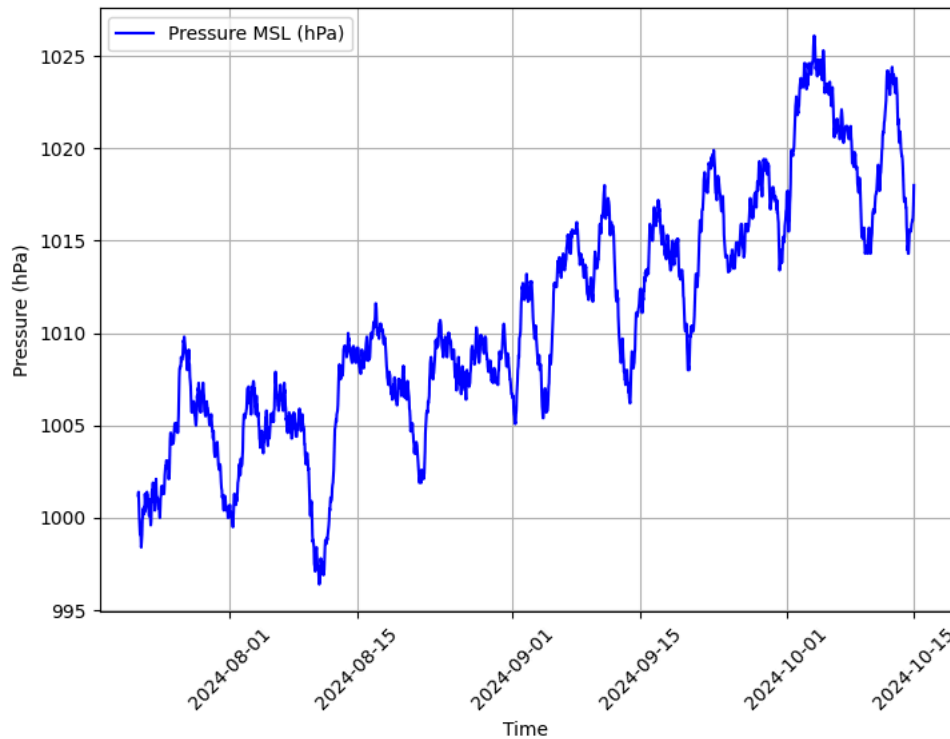
# Plotting Mean Sea Level Pressure
plt.subplot(1, 2, 1)
plt.plot(hourly_dataframe['date'], hourly_dataframe['surface_pressure'], color='blue', label='Pressure MSL (hPa)')
plt.title('Mean Sea Level Pressure on 2023-10-01')
plt.xlabel('Time')
plt.ylabel('Pressure (hPa)')
plt.xticks(rotation=45)
plt.grid()
plt.legend()

```

```
# Adjust layout
plt.tight_layout()
plt.show()
```

```
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.54.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.7)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.26.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
```

Mean Sea Level Pressure on 2023-10-01




```
import numpy as np
NOISE_LEVEL = 1.0 # Example noise level
NUM_INIT_POINTS = 2040 # Example number of initial points

# Generate random samples from a normal distribution
random_values = np.random.normal(0., NOISE_LEVEL, NUM_INIT_POINTS)

# Create a Pandas DataFrame
df = pd.DataFrame(random_values, columns=['surface_pressure'])
x1=pd.DataFrame(hourly_dataframe['date'])
y=pd.DataFrame(hourly_dataframe['surface_pressure'])
yy=df+y


plt.figure(dpi=100)
plt.xlabel("$x$")
plt.ylabel("$y$")
plt.scatter(x1, y, marker='o', c='k', zorder=1, s=.25, label='No noise')
plt.ylim(995, 1030);
```

Enable brow
Settings to
executions



Enable browser
Settings to
allow automatic
executions

```
plt.figure(dpi=100)
plt.xlabel("$x$")
plt.ylabel("$y$")
plt.scatter(x1[1:50], yy[1:50], marker='o', c='k', zorder=1, s=5, label='Noisy observations')
plt.ylim(995, 1005);
```



```
import gpax
import numpy as np
import matplotlib.pyplot as plt

gpax.utils.enable_x64() # enable double precision

yy=yy.to_numpy()

x1 = pd.DataFrame(hourly_dataframe['date'])
x1['unix_time'] = hourly_dataframe['date'].astype('int64') // 10**9 # Convert to seconds
x1['unix_time'] = x1['unix_time'] / 10**9

x1=x1.to_numpy()

x1=x1[:,1]

xx1=jnp.array(x1.astype(np.float32))
```

```
print(xx1)
```

```
→ [1.7216064 1.72161 1.7216136 ... 1.7289397 1.7289432 1.7289468]
```

Standard GP

Next, we initialize and train a GP model. We are going to use an RBF kernel, $k_{RBF} = \exp(-\frac{\|x_i - x_j\|^2}{2l^2})$, which is a "go-to" kernel functions in GP.

```
import jax.numpy as jnp
# Get random number generator keys for training and prediction
key1, key2 = gpax.utils.get_keys()
```

```
gp_model.fit(key1, jnp.array(xx1[:,2]), jnp.array(yy[:,2]), num_chains=1, num_warmup=20, num_samples=20)
```

```
→ sample: 100%|██████████| 40/40 [02:15<00:00, 3.38s/it, 7 steps of size 6.81e-02. acc. prob=0.87]
      mean      std      median      5.0%      95.0%      n_eff      r_hat
k_length[0]    0.07    0.09    0.00    0.00    0.22    3.68    1.61
k_scale  77646.89 23059.94 76252.73 37432.96 102534.61 12.70    0.99
noise      9.75    0.78    9.45    8.66    11.03    3.77    1.72
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-97-de6fb658a02e> in <cell line: 8>()
      6 gp_model.fit(key1, jnp.array(xx1[:,2]), jnp.array(yy[:,2]), num_chains=1, num_warmup=20, num_samples=20)
      7
----> 8 current_loss = gp_model.loss() # Replace with your method to compute loss

AttributeError: 'ExactGP' object has no attribute 'loss'
```

Next steps: [Explain error](#)

"Now let's use our trained model to obtain a probabilistic prediction of function values on new ("test") data. In the fully Bayesian mode, we get a pair of predictive mean and covariance, μ_*^i and Σ_*^i , for each i -th HMC sample with kernel parameters θ ,

$$\mu_*^i = K(X_*, X|\theta^i)K(X, X|\theta^i)^{-1}y$$

$$\Sigma_*^i = K(X_*, X_*|\theta^i) - K(X_*, X|\theta^i)K(X, X|\theta^i)^{-1}K(X, X_*|\theta^i)$$

The `.predict()` method returns the center of the mass of all the predictive means,

$$\mu_*^{post} = \frac{1}{L} \sum_{i=1}^L \mu_*^i,$$

which corresponds to the `posterior_mean` in the code cell below, and samples

$$f_*^i \sim MVNormal(\mu_*^i, \Sigma_*^i)$$

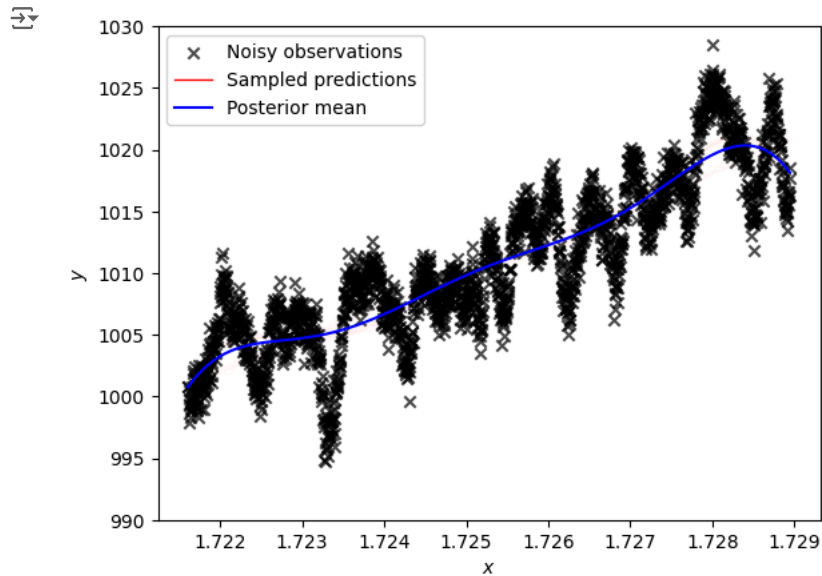
from multivariate normal distributions for all the pairs of predictive means and covariances (`f_samples` in the code cell below). Note that model noise is absorbed into the kernel computation function."

```
# Prepare test inputs
X_test = np.linspace(1.7216064, 1.7289468, 100)
# Get the GP prediction. Here n stands for the number of samples from each MVNormal distribution
# (the total number of MVNormal distributions is equal to the number of HMC samples)
posterior_mean, f_samples = gp_model.predict(key2, X_test, n=200)
```

Plot the obtained results:

```
_, ax = plt.subplots(dpi=100)
ax.set_xlabel("$x$")
ax.set_ylabel("$y$")
ax.scatter(jnp.array(xx1), jnp.array(yy), marker='x', c='k', zorder=1, label="Noisy observations", alpha=0.7)
for y1 in f_samples:
    ax.plot(X_test, y1.mean(0), lw=1, zorder=0, c='r', alpha=1)
l, = ax.plot(X_test, f_samples[0].mean(0), lw=1, c='r', alpha=1, label="Sampled predictions")
ax.plot(X_test, posterior_mean, lw=1.5, zorder=1, c='b', label='Posterior mean')
```

```
ax.legend(loc='upper left')
l.set_alpha(0)
ax.set_ylim(990, 1030);
```



Enable browser
Settings to
executions

We can see there is a relatively large uncertainty - defined by the dispersion in sampled predictions - in GP prediction between -0.75 and -0.25 where no measurements are available. It is also common in the GP literature to draw the GP uncertainty as a 2σ shaded region around the GP posterior mean:

```
_, ax = plt.subplots(dpi=100)
ax.set_xlabel("$x$")
ax.set_ylabel("$y$")
ax.scatter(jnp.array(xx1), jnp.array(yy), marker='x', c='k', zorder=2, label="Noisy observations", alpha=0.7)
ax.plot(X_test, posterior_mean, lw=1.5, zorder=2, c='b', label='Posterior mean')
ax.fill_between(X_test,
                posterior_mean - f_samples.std(axis=(0,1)),
                posterior_mean + f_samples.std(axis=(0,1)),
                color='r', alpha=0.3, label="Model uncertainty")
ax.legend(loc='upper left')
ax.set_ylim(990, 1030);
```

